

Concurrency Control in Distributed Systems

* PROBLEM STATEMENT

User 1 →

User 2 →

User 3 →

Seat	
id	status
10	FREE

{
Read Row seat with id: 10
if FREE:
Logic to change state from FREE to BOOKED
Update the DB.
}

CRITICAL SECTION
(in application logic,
shared resource)

Solution.

① Using 'SYNCHRONISED' for critical section.

U1 ✓
U2 X
U3 X

} won't work for distributed
concurrency.

1 Process → n threads ✓ synchronised work

n Process X synchronised won't work

(n machines)

② Using DISTRIBUTED CONCURRENCY CONTROL

OPTIMISTIC (OCC)

PESSIMISTIC (PCC)

Recap:-

① Usage of TRANSACTION?

② DB LOCKING?

③ ISOLATION LEVEL present ?

we apply
at DB level.

we apply at
application
code
level

→ helps achieve INTEGRITY.

i.e. helps avoid INCONSISTENCY in DB.

tx. db. begin()

- debit money from A

- credit money to B

if all success

commit();

else rollback();

id	user	amount	type
1	A		Credit
2	B		debit



DB

tx.db.end()

rollback()

TODO: know more about these work internally.

→ if tm-fails, it will rollback all nested statements run as part of this txn.

DB LOCKING

LOCK
SHARED:(S)

EXCLUSIVE:(X)

→ makes sure that no other txn updates the locked rows.

LOCK TYPE	ANOTHER SHARED LOCK?	ANOTHER X LOCK?
HAVE SHARED LOCK	✓	X
HAVE (X) LOCK	X	X

(S):- is used for reads.

- in a shared lock, only read can happen

- restricts any other txn. to update record on which S is acquired

T₁ (S) read [13 | Free | .. |]
T₂ read ✓

(X):-

- is used by a txn. to WRITE to a record
/ update

update

- other txns. cannot read/write to a record on which (X) is taken.


→ If a row has shared lock, first remove shared lock, then grant exclusive lock.

→ On exclusive lock, no other txn is allowed to take shared/exclusive lock.

ISOLATION LEVEL

- how much level of concurrency is allowed in your DATABASE.

- ACID
↓
ISOLATION



these txns feel like they are working alone.

PROBLEMS

ISOLATION LEVEL	^{causes problems} DIRTY READ	NON-REPEATABLE READ	PHANTOM READ
1) READ UNCOMMITTED	Yes	Yes	Yes
2) READ COMMITTED	No	Yes	Yes
3) REPEATABLE READ	No	No	Yes

③ REPEATABLE READ

④ SERIALIZABLE

No

No

No

① DIRTY READ PROBLEM :-

- if txn A is reading data written by txn B but not yet committed.
- if txn B does rollback(), then whatever data read by txn A is dirty read.

② NON-REPEATABLE READ :-

- if txn A reads same data/record multiple times, in a txn & some data does not get.

↳ other txn B has successfully committed for that record.
commit()

③ PHANTOM READ :-

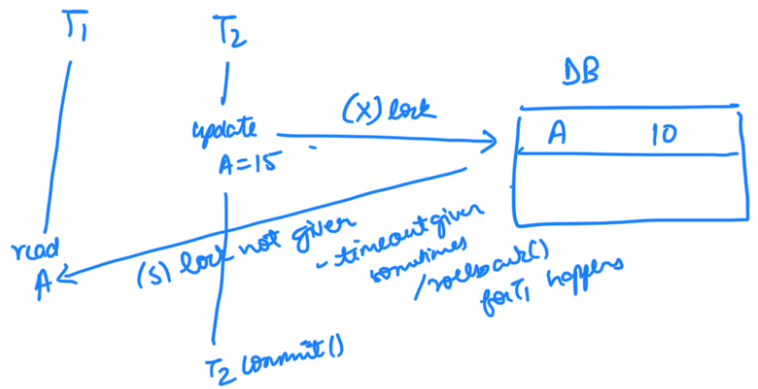
- txn A executes same query several times, & rows returned are different.
value ≥ 10 & value ≤ 15

within range


txn B commits new rows to the database.

First read this, then go to problems table

ISOLATION LEVEL	CONCURRENCY HIGH	LOCKING STRATEGY (S/X)
① READ UNCOMMITTED		Read: No lock Write: No lock (Very risky) (When you require only reading)
② READ COMMITTED		Read: Acquire S, Release as read is done. Write: Acquire X, & keep till end of txn. i.e. end/commit()



\therefore Dirty Read not possible



\therefore Non-Repeatable Read possible

Phantom Read possible

LESS CONCURRENT THAN READ

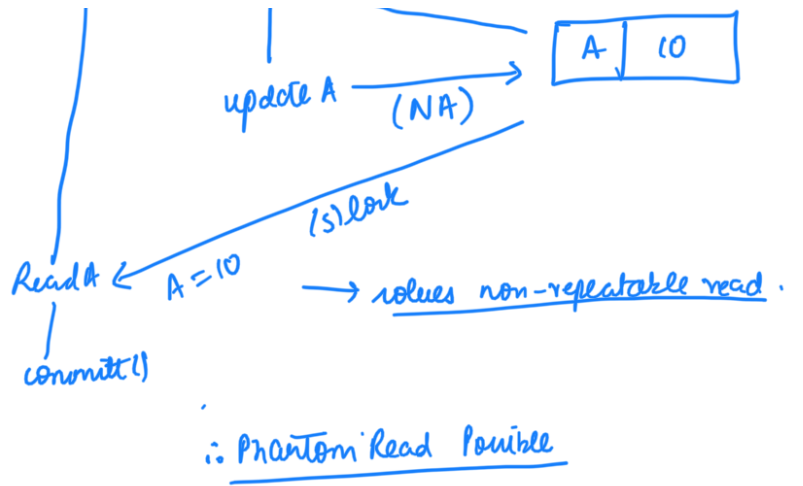
UNCOMMITTED

③ REPEATABLE READ

Read:- (S) lock acquired, release only at end of txn.

Write:- (X) lock acquired, release only at end of txn.

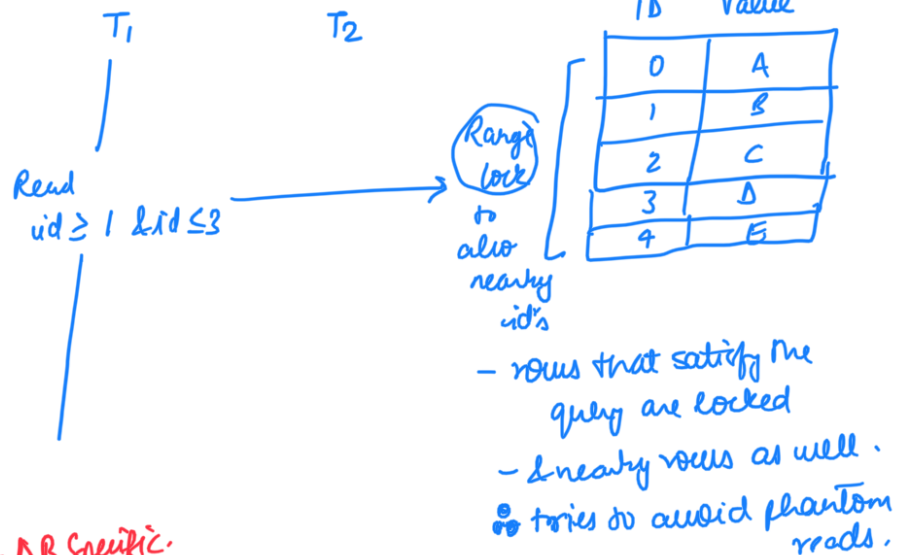




④ SERIALIZABLE

REPEATABLE READ STRATEGY
+
RANGE LOCK
(lock is released at the end of txn)

CONCURRENCY LOW



we set id for every transaction
else, default is selected for DB specific.

SET TRANSACTION ISOLATION LEVEL REPEATABLE READ;

BEGIN TRANSACTION

select

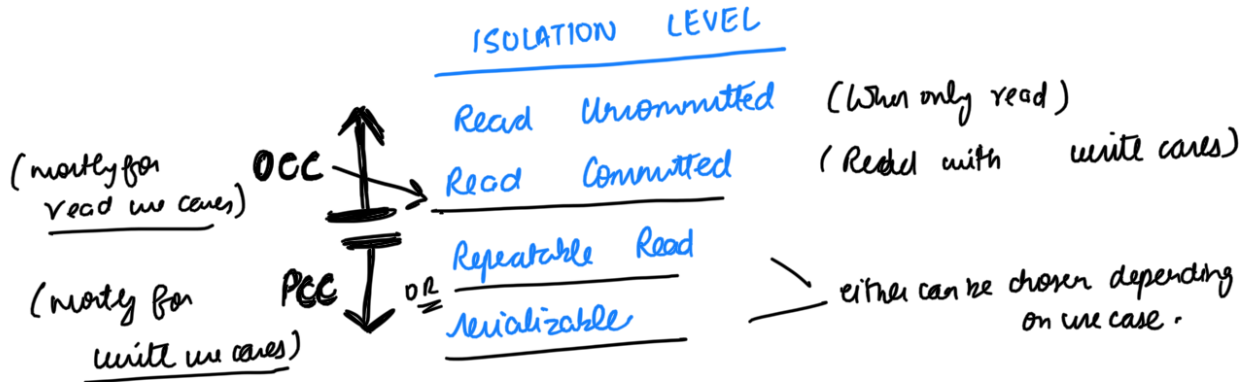
update

if success: COMMIT TRANSACTION

else: ROLLBACK TRANSACTION

END TRANSACTION

OCC & PCC



OCC

(Read Committed ISOLATION LEVEL)

- solves concurrency using version.

↳ additional input row in DB

concurrent txns.

Txn A

Txn B

DB

Begin TRANSACTION

Begin Transaction

ID:1

Status: Free

Version: 1

Read Row id:1

Read Row id:1

93

Version:1

(S) lock with release.

Version:1

(S) lock with release

select for update

(Version Validation Happens)

matching versions

Version:1

(X) lock by txn A

update Row ID:1

ID:1

Status: Locked

Time



Status: Booked
Version: 2

has acquired (X)
txn B cannot do
read / write till
commit / abort

Status: Booked
Version: 2

(X) lock by txn A

Commit
↳ Releases
exclusion
lock

ID: 1
Status: Booked
Version: 2

Select for Update
(Version Validation)

Version: 1

11

Version: 2

(X) lock applied.

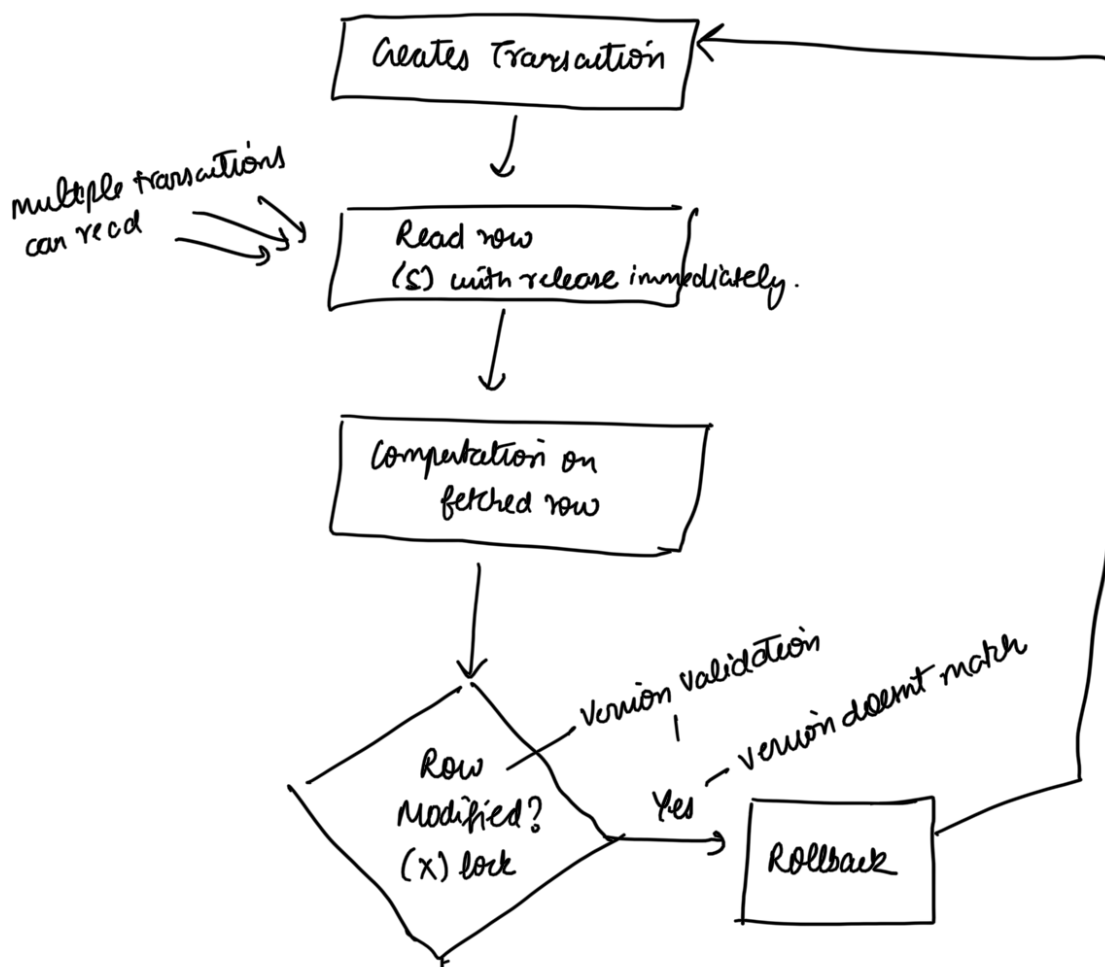
Version mismatch

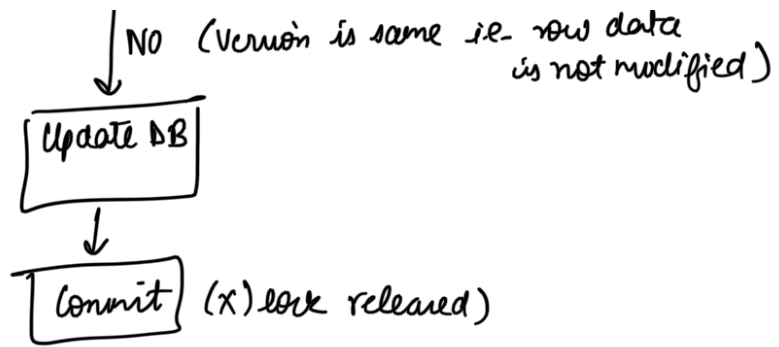
Rollback()

19

(X) lock released

Flow for a Transaction :-





Mostly used in companies :-

- very less sudden concurrency comes, OCC will handle.

PCC

ISOLATION LEVEL

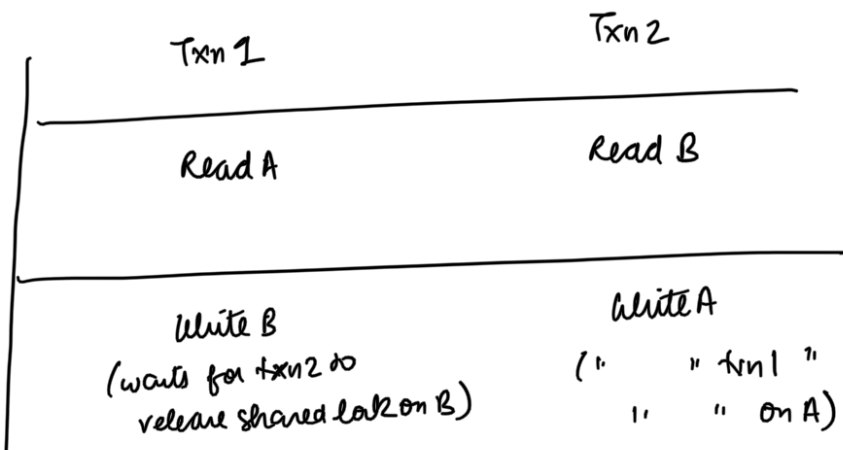
- REPEATABLE READ
- SERIALIZABLE

- brings serialization.


(sequence because of locking)

- other txns would have to wait to update before one txn is done committing the update.

- this creates a problem of deadlocks in PCC



DB		
idn	value	
1	A	→ (S) lock by Txn1
2	B	→ (S) lock by Txn2



deadlock

- waiting for INF
- about txns
- all locks released
- start new transactions.

∴ OCC solves deadlock over PCC

OCC

* ISOLATION LEVEL

Read:- Read Uncommitted

✓ Write:- Read Committed

* High Concurrency

* No chance of deadlock

* In case of conflict,

✓ overhead of transaction rollback (check versions)

& retry logic is there

PCC

* Repeatable Read & Serializable

* Low concurrency compared to optimistic

✓ * Deadlock is possible, txns stuck in deadlock are forced to rollback.

✓ * Putting a long lock (isolation level duration) in case of long running txn

↓
timeout time happens & rollback needs to be done.